

Loading Necessary Modules

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4
5 from sklearn import model_selection
6
7 from keras.wrappers.scikit_learn import KerasRegressor
8 from keras.models import Sequential
9 from keras.layers import Dense
10 from keras.layers import LSTM
11 from keras.layers import BatchNormalization
12 from keras.layers import Dropout
13 from keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau, LearningRateScheduler
14
15 import tensorflow as tf
16
17 import helper
```

Using TensorFlow backend.

Styling Tables

In [2]:

```
1 %%HTML
2 <style type='text/css'>
3 table.dataframe th, table.dataframe td{
4     border: 3px solid purple !important;
5     color: solid black !important;
6 }
7 </style>
```

Loading Data

In [3]:

```
1 # Loading dataset
2 filename = "Clean_Akosombo_data.csv"
3 akosombo = helper.load_csv_data(filename)
```

Successfully loaded!

Splitting Data

In [4]:

```
1 # Splitting dataset
2 target_variable = "generation"
3 X, y, X_train, X_test, X_val, y_train, y_test, y_val = helper.split_data(akosombo, target_variable)
```

Data is splitted into X, y, X_train, X_test, X_val, y_train, y_test, y_val.

Shape Info of Features Training Set:

Number of datapoints (rows): 7500

Number of features (columns): 2

Shape Info of Features Test Set:

Number of datapoints (rows): 2501

Number of features (columns): 2

Shape Info of Features Validation Set:

Number of datapoints (rows): 2501

Number of features (columns): 2

Scaling Data

In [5]:

```
1 # Data Scaling
2 X_train, X_test, X_val = helper.scale(X_train, X_test, X_val, scale_validation=True)
```

Model Creation

In [9]:

```

1  # Creating Sequential Model
2  neural_network_model = Sequential()
3
4  # Input Layer
5  neural_network_model.add(Dense(20, input_dim=X_train.shape[1], kernel_initializer='normal'))
6
7  # Hidden Layers
8  neural_network_model.add(Dense(40, kernel_initializer='normal', activation='relu'))
9  neural_network_model.add(Dense(40, kernel_initializer='normal', activation='relu'))
10 neural_network_model.add(Dropout(0.01))
11
12 # Output Layer
13 neural_network_model.add(Dense(1, kernel_initializer='normal', activation='linear'))
14
15 # Compiling the network
16 neural_network_model.compile(loss='mean_squared_error', optimizer='adam', metrics=['mean_squared_error'])
17 neural_network_model.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 20)	60
dense_6 (Dense)	(None, 40)	840
dense_7 (Dense)	(None, 40)	1640
dropout_2 (Dropout)	(None, 40)	0
dense_8 (Dense)	(None, 1)	41
Total params: 2,581		
Trainable params: 2,581		
Non-trainable params: 0		

Callbacks

In [10]:

```
1 checkpoint_name = 'Weights-{epoch:03d}--{val_loss:.5f}.h5'
2
3
4 checkpoint = ModelCheckpoint(checkpoint_name,
5                             monitor="val_loss",
6                             save_best_only=True,
7                             verbose=1,
8                             mode='auto')
9
10
11 earllystop = EarlyStopping(monitor='val_loss',
12                             min_delta=0,
13                             patience=3,
14                             verbose=1,
15                             mode='auto',
16                             restore_best_weights=True)
17
18
19 reduce_lr = ReduceLROnPlateau(monitor='val_loss',
20                                factor=0.2,
21                                patience=3,
22                                verbose=1,
23                                mode='auto',
24                                min_delta=0.00001)
25
26 # Putting call backs into a callback list
27 callbacks = [earllystop, checkpoint, reduce_lr]
```

Training Model

In [11]:

```

1 epochs = 150
2 batch_size = 5
3
4 history = neural_network_model.fit(
5     X_train, y_train,
6     batch_size=batch_size,
7     validation_data=(X_test, y_test), # Overrides validation_split argument.
8     # validation_split=0.25,
9     epochs=epochs,
10    verbose=2,
11    callbacks=callbacks,
12 )

```

Train on 7500 samples, validate on 2501 samples

Epoch 1/150

- 2s - loss: 9.7737 - mean_squared_error: 9.7738 - mean_absolute_error: 1.2441 - val_loss: 0.1738 - val_mean_squared_error: 0.1738 - val_mean_absolute_error: 0.3295

Epoch 00001: val_loss improved from inf to 0.17377, saving model to Weights-001--0.17377.h5

Epoch 2/150

- 1s - loss: 0.2522 - mean_squared_error: 0.2522 - mean_absolute_error: 0.3883 - val_loss: 0.1546 - val_mean_squared_error: 0.1546 - val_mean_absolute_error: 0.2972

Epoch 00002: val_loss improved from 0.17377 to 0.15459, saving model to Weights-002--0.15459.h5

Epoch 3/150

- 2s - loss: 0.2640 - mean_squared_error: 0.2640 - mean_absolute_error: 0.3929 - val_loss: 0.1541 - val_mean_squared_error: 0.1541 - val_mean_absolute_error: 0.2962

Epoch 00003: val_loss improved from 0.15459 to 0.15405, saving model to Weights-003--0.15405.h5

Epoch 4/150

- 1s - loss: 0.2759 - mean_squared_error: 0.2759 - mean_absolute_error: 0.3996 - val_loss: 0.1708 - val_mean_squared_error: 0.1708 - val_mean_absolute_error: 0.3255

Epoch 00004: val_loss did not improve from 0.15405

Epoch 5/150

- 1s - loss: 0.2613 - mean_squared_error: 0.2613 - mean_absolute_error: 0.3945 - val_loss: 0.1462 - val_mean_squared_error: 0.1462 - val_mean_absolute_error: 0.2946

Epoch 00005: val_loss improved from 0.15405 to 0.14619, saving model to Weights-005--0.14619.h5

Epoch 6/150

- 1s - loss: 0.2584 - mean_squared_error: 0.2584 - mean_absolute_error: 0.3859 - val_loss: 0.1515 - val_mean_squared_error: 0.1515 - val_mean_absolute_error: 0.2979

Epoch 00006: val_loss did not improve from 0.14619

Epoch 7/150

- 1s - loss: 0.2597 - mean_squared_error: 0.2597 - mean_absolute_error: 0.3920 - val_loss: 0.2564 - val_mean_squared_error: 0.2564 - val_mean_absolute_error: 0.4028

Epoch 00007: val_loss did not improve from 0.14619

Epoch 8/150

- 1s - loss: 0.2711 - mean_squared_error: 0.2711 - mean_absolute_error: 0.3987 - val_loss: 0.1469 - val_mean_squared_error: 0.1469 - val_mean_absolute_error: 0.2909

Restoring model weights from the end of the best epoch

Epoch 00008: val_loss did not improve from 0.14619

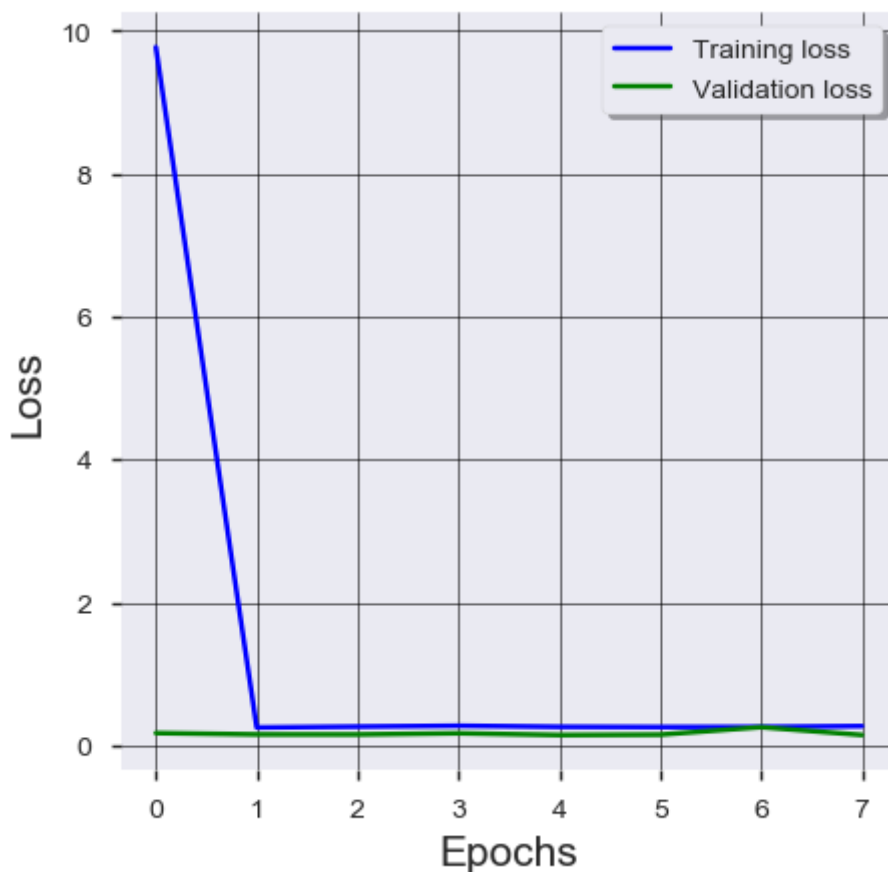
Epoch 00008: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 00008: early stopping

Visualizing Losses

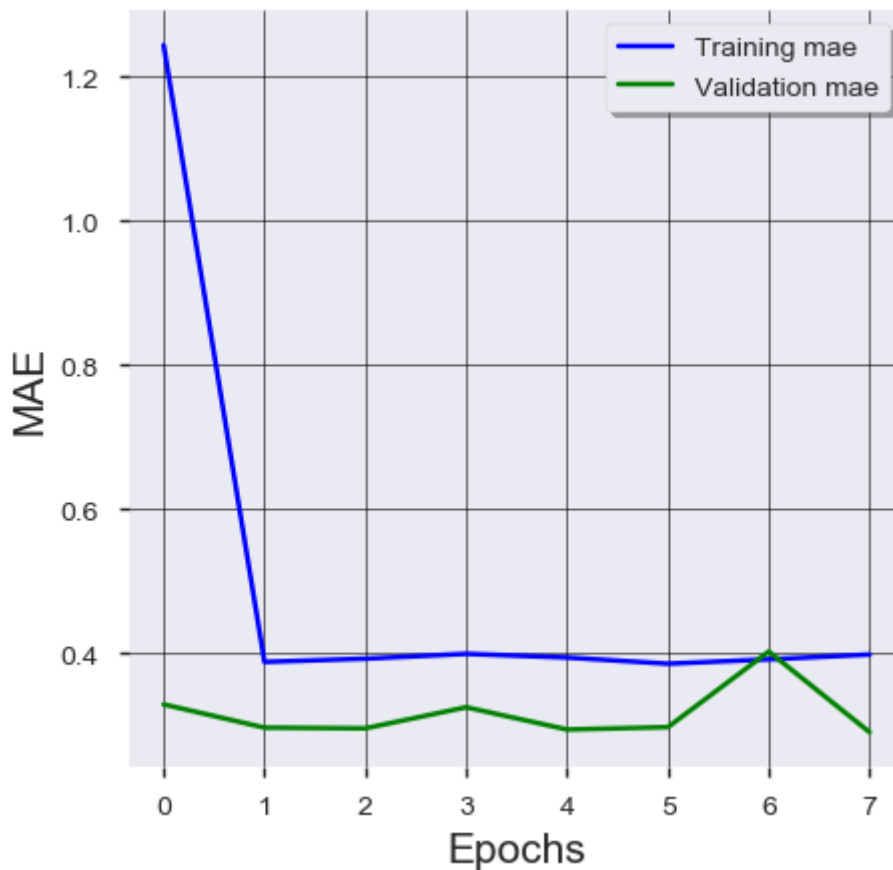
In [12]:

```
1 training_loss = history.history['loss']
2 validation_loss = history.history['val_loss']
3 epochs = history.epoch
4
5 fig = plt.figure(figsize=(5,5), dpi=100)
6
7 plt.plot(epochs, training_loss, label='Training loss', color='blue')
8 plt.plot(epochs, validation_loss, label='Validation loss', color='green')
9
10 plt.xlabel('Epochs', fontsize=15)
11 plt.ylabel('Loss', fontsize=15)
12
13 plt.grid(b=True, which="both", axis="both", color="black", linewidth=0.4)
14 plt.legend(loc='best', fontsize='medium', numpoints=1, frameon=True, shadow=True, fancy)
15
16 plt.savefig('Neural_Network_Loss.png', dpi=300, transparent=True)
17
18 plt.show()
```



In [13]:

```
1 training_accuracy = history.history['mean_absolute_error']
2 validation_loss = history.history['val_mean_absolute_error']
3 epochs = history.epoch
4
5
6 fig = plt.figure(figsize=(5,5), dpi=100)
7
8 plt.plot(epochs, training_accuracy, label='Training mae', color='blue')
9 plt.plot(epochs, validation_loss, label='Validation mae', color='green')
10
11 plt.xlabel('Epochs', fontsize=15)
12 plt.ylabel('MAE', fontsize=15)
13
14 plt.grid(b=True, which="both", axis="both", color="black", linewidth=0.4)
15 plt.legend(loc='best', fontsize='medium', numpoints=1, frameon=True, shadow=True, fancy
16
17 plt.savefig('Neural_Network_mean_absolute_error.png', dpi=300, transparent=True)
18
19 plt.show()
```



Model History Data

In [39]:

```
1 history_data = pd.DataFrame(history.history)
2 history_data['epochs'] = history.epoch
3 history_data.to_csv("neural_network_history.csv", index=True)
4 history_data
```

Out[39]:

	val_loss	val_mean_squared_error	val_mean_absolute_error	loss	mean_squared_error
0	0.173769	0.173769	0.329457	9.773736	9.773754
1	0.154588	0.154588	0.297226	0.252188	0.252188
2	0.154051	0.154051	0.296220	0.263980	0.263980
3	0.170786	0.170786	0.325498	0.275920	0.275920
4	0.146188	0.146188	0.294601	0.261313	0.261313
5	0.151516	0.151516	0.297905	0.258412	0.258412
6	0.256357	0.256357	0.402799	0.259739	0.259739
7	0.146919	0.146919	0.290868	0.271056	0.271056

In [40]:

```
1 model_parameters = pd.DataFrame(history.params)
2 model_parameters.to_csv("neural_network_parameters.csv", index=True)
3 model_parameters
```

Out[40]:

	batch_size	epochs	steps	samples	verbose	do_validation	metrics
0	5	150	None	7500	2	True	loss
1	5	150	None	7500	2	True	mean_squared_error
2	5	150	None	7500	2	True	mean_absolute_error
3	5	150	None	7500	2	True	val_loss
4	5	150	None	7500	2	True	val_mean_squared_error
5	5	150	None	7500	2	True	val_mean_absolute_error

Model Evaluation


In [16]:

```

1 model_name = 'Neural Network'
2 helper.evaluate(X_test, y_test, model_name, neural_network_model)

```

Neural Network Mean Squared Error: 0.14618774985909064
 Neural Network Root Mean Squared Error: 0.38234506647672417
 Neural Network R2 Score: 0.9697627857246928
 Neural Network Explained Variance Score: 0.9698258542620452
 Neural Network Mean Absolute Error: 0.29460089848033
 Neural Network Median Absolute Error: 0.2666202926635748
 Neural Network Mean Squared Log Error: 0.0007157302181291301



Prediction

In [17]:

```

1 y_pred = neural_network_model.predict(X_test)

```

In [18]:

```

1 model_prediction_results = pd.DataFrame({
2     'actual_generation' : list(y_test),
3     'predicted_generation' : list(y_pred),
4 })

```

In [19]:

```
1 model_prediction_results.head()
```

Out[19]:

	actual_generation	predicted_generation
0	17.596	[17.10944]
1	15.630	[15.234886]
2	10.850	[11.42171]
3	14.520	[14.833691]
4	8.420	[8.61916]

In [20]:

```
1 model_prediction_results.dtypes
```

Out[20]:

```
actual_generation      float64
predicted_generation    object
dtype: object
```

In [21]:

```
1 # data['y_pred_generation'].str.replace('[\[\]]', '')
```

In [22]:

```
1 model_prediction_results = model_prediction_results.astype({'predicted_generation': 'flo
```

In [23]:

```
1 model_prediction_results.head(10)
```

Out[23]:

	actual_generation	predicted_generation
0	17.596	17.109440
1	15.630	15.234886
2	10.850	11.421710
3	14.520	14.833691
4	8.420	8.619160
5	13.640	14.780686
6	14.010	13.918062
7	9.699	9.447412
8	13.900	12.787080
9	11.150	11.148086

In [24]:

```
1 model_prediction_results.to_csv("neural_network_predicted_values.csv", index=True)
```

Saving the Model

In [25]:

```
1 #Saving model
2 neural_network_model.save('neural_network_model.h5')
```

Optimizing the Hyperparameter of the Neural Network with GridSearchCV

In [26]:

```
1 # Setting Random Seed
2 np.random.seed(82)
```

In [27]:

```
1 def neural_network_regressor_model(optimizer, activation):
2     model = Sequential()
3
4     # Input Layer
5     model.add(Dense(20, input_dim=X_train.shape[1], kernel_initializer='normal', activation=activation))
6
7     # Hidden Layers
8     model.add(Dense(40, kernel_initializer='normal', activation=activation))
9     model.add(Dense(40, kernel_initializer='normal', activation=activation))
10
11    # Output Layer
12    model.add(Dense(1, kernel_initializer='normal', activation='linear'))
13
14    # Compiling the network :
15    model.compile(loss='mean_squared_error', optimizer=optimizer, metrics=['mean_squared_error'])
16
17    print(model.summary())
18    return model
```

In [28]:

```
1 epochs = 150
2 batch_size = 5
3
4 nn_model = KerasRegressor(
5     build_fn=neural_network_regressor_model,
6     epochs=epochs,
7     batch_size=batch_size,
8     verbose=2
9 )
```

In []:

1

In [29]:

```

1  # Kfold with with n_splits = 5 to split the Dataset into 5-folds
2  kfold = model_selection.KFold(n_splits=5, shuffle=True, random_state=23)
3
4  # Dictionary of parameters to optimize
5  parameters = {
6      "activation" : ['tanh', 'relu', 'elu', 'selu'],
7      "optimizer" : ['adam', 'rmsprop'],
8      "batch_size" : [5, 10, 15, 20],
9      "epochs" : [50, 100, 150],
10 }
11
12 # Scoring Metric
13 scorer = "r2"
14
15 # Instantiating Search object
16 grid = model_selection.RandomizedSearchCV(
17     estimator=nn_model,
18     param_distributions=parameters,
19     scoring=scorer,
20     cv=kfold,
21     n_jobs=1,
22     verbose=2,
23 )
24
25 # Fit the grid object on Training Dataset
26 grid.fit(X_train, y_train)

```

```

- 0s - loss: 0.1532 - mean_squared_error: 0.1532 - mean_absolute_error:
0.3095
Epoch 134/150
- 0s - loss: 0.1537 - mean_squared_error: 0.1537 - mean_absolute_error:
0.3112
Epoch 135/150
- 0s - loss: 0.1553 - mean_squared_error: 0.1553 - mean_absolute_error:
0.3103
Epoch 136/150
- 0s - loss: 0.1533 - mean_squared_error: 0.1533 - mean_absolute_error:
0.3101
Epoch 137/150
- 0s - loss: 0.1531 - mean_squared_error: 0.1531 - mean_absolute_error:
0.3098
Epoch 138/150
- 0s - loss: 0.1541 - mean_squared_error: 0.1541 - mean_absolute_error:
0.3105
Epoch 139/150
- 0s - loss: 0.1528 - mean_squared_error: 0.1528 - mean_absolute_error:
0.3079

```

In [30]:

```

1 # Saving Hyperparameter optimization results as a DataFrame
2 results = pd.DataFrame(grid.cv_results_[["params", "mean_test_score", "std_test_score", "rank_test_score"]])
3 results.sort_values("rank_test_score", inplace=True)
4 results.to_csv("neural_network_hyperparameter_optimization_results.csv", index=True)
5 results

```

Out[30]:

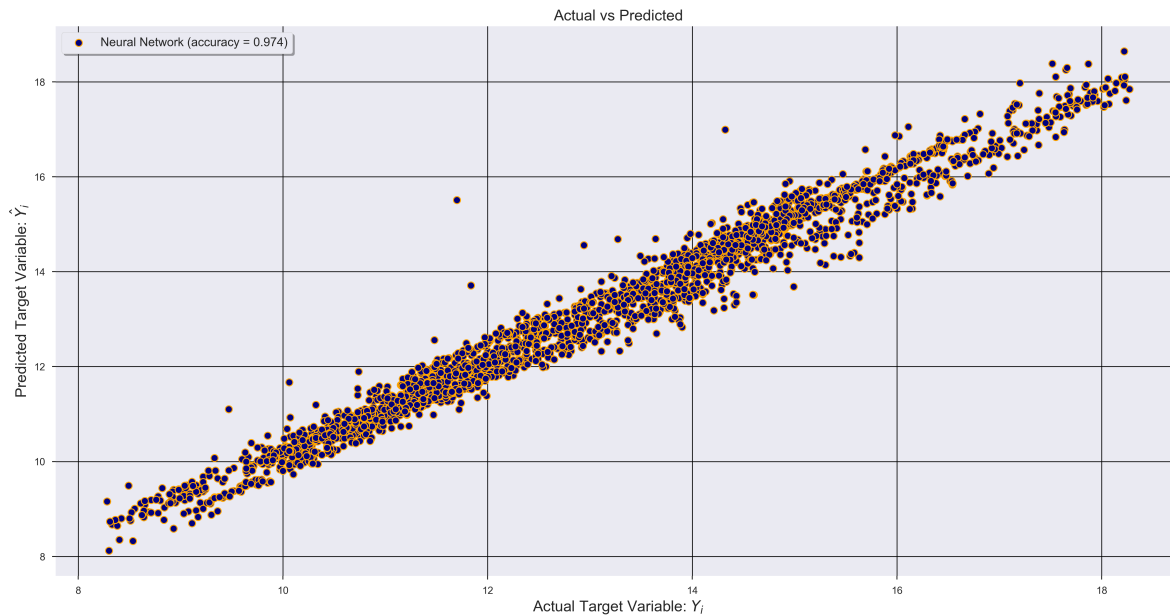
	params	mean_test_score	std_test_score	rank_test_score
2	{'optimizer': 'adam', 'epochs': 50, 'batch_size': 128}	0.975144	0.001759	1
4	{'optimizer': 'adam', 'epochs': 150, 'batch_size': 128}	0.975027	0.002148	2
5	{'optimizer': 'adam', 'epochs': 50, 'batch_size': 64}	0.973465	0.001229	3
6	{'optimizer': 'adam', 'epochs': 100, 'batch_size': 128}	0.972549	0.001911	4
3	{'optimizer': 'adam', 'epochs': 100, 'batch_size': 64}	0.972254	0.001825	5
8	{'optimizer': 'rmsprop', 'epochs': 150, 'batch_size': 128}	0.971457	0.001375	6
7	{'optimizer': 'adam', 'epochs': 100, 'batch_size': 64}	0.969546	0.004080	7
9	{'optimizer': 'rmsprop', 'epochs': 50, 'batch_size': 128}	0.968880	0.007177	8
1	{'optimizer': 'rmsprop', 'epochs': 50, 'batch_size': 64}	0.968604	0.004121	9
0	{'optimizer': 'adam', 'epochs': 50, 'batch_size': 128}	0.966260	0.003683	10

Evaluating Best Estimator

In [31]:

```
1 best_estimator = grid.best_estimator_  
2 helper.evaluate(X_test, y_test, "Neural Network", best_estimator)
```

Neural Network Mean Squared Error: 0.1242707793972232
Neural Network Root Mean Squared Error: 0.3525206084716512
Neural Network R2 Score: 0.9742960529290915
Neural Network Explained Variance Score: 0.9749624975053861
Neural Network Mean Absolute Error: 0.27229552823035874
Neural Network Median Absolute Error: 0.22310390472412145
Neural Network Mean Squared Log Error: 0.0006259880454793604



Predicting with the Best Estimator and Saving Predicted Results as csv

In [34]:

```

1 tune_y_pred = best_estimator.predict(X_test)
2
3 hyp_tune_data = pd.DataFrame(
4     {"actual_generation": list(y_test), "predicted_generation": list(tune_y_pred),}
5 )
6
7 hyp_tune_data.to_csv("best_estimator_predicted_values.csv", index=True)
8
9 hyp_tune_data.head(10)

```

Out[34]:

	actual_generation	predicted_generation
0	17.596	17.467146
1	15.630	15.372869
2	10.850	11.452371
3	14.520	14.715870
4	8.420	8.799006
5	13.640	14.690557
6	14.010	13.920854
7	9.699	9.571536
8	13.900	12.830315
9	11.150	11.287091

In []:

1

Saving the Model

In [35]:

```

1 import joblib
2
3 joblib.dump(best_estimator, "optimized_neural_network_model.joblib")

```

Out[35]:

```
['optimized_neural_network_model.joblib']
```

In []:

1

In [38]:

```
1 best_estimator.get_params()
```

Out[38]:

```
{'epochs': 50,  
 'batch_size': 5,  
 'verbose': 2,  
 'optimizer': 'adam',  
 'activation': 'selu',  
 'build_fn': <function __main__.neural_network_regressor_model(optimizer, ac  
tivation)>}
```

In []:

```
1
```