

Loading all necessary Packages/Libraries

In [1]:

```

1 # Loading the iconic trio 🎶
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 # Importing model_selection to get access to some dope functions Like GridSearchCV()
7 from sklearn import model_selection
8
9 # from sklearn.externals import joblib
10
11 # Loading models
12 from sklearn import linear_model
13 from sklearn import svm
14 from sklearn import tree
15 from sklearn import ensemble
16 from sklearn import neighbors
17
18 # custom
19 import helper
20
21 # Loading black for formatting codes
22 %load_ext blackcellmagic

```

Styling Tables

In [2]:

```

1 %%HTML
2 <style type='text/css'>
3 table.dataframe th, table.dataframe td{
4     border: 3px solid purple !important;
5     color: solid black !important;
6 }
7 </style>

```

Loading the Dataset

In [3]:

```

1 # Loading dataset
2 filename = "Clean_Akosombo_data.csv"
3 akosombo = helper.load_csv_data(filename)

```

Successfully loaded!

Splitting the Dataset

In [4]:

```
1 # Splitting dataset
2 target_variable = "generation"
3 X, y, X_train, X_test, y_train, y_test = helper.split_data(akosombo, target_variable)
```

Data is splitted into X, y, X_train, X_test, y_train, y_test.

Shape Info of Features Training Set:

Number of datapoints (rows): 10001

Number of features (columns): 2

Shape Info of Features Test Set:

Number of datapoints (rows): 2501

Number of features (columns): 2

Scaling the Dataset

In [5]:

```
1 # Data Scaling
2 X_train, X_test = helper.scale(X_train, X_test)
```

Choosing Baseline Models and Training Models

In [6]:

```

1 # Instantiating baseline models
2 models = [
3     ("Linear Regression", linear_model.LinearRegression()),
4     # ("Lasso", linear_model.Lasso()),
5     ("Ridge", linear_model.Ridge()),
6     # ("SGD", linear_model.SGDRegressor()),
7     ("SVR", svm.LinearSVR()),
8     ("NuSVR", svm.NuSVR()),
9     ("SVR", svm.SVR()),
10    ("Decision Tree", tree.DecisionTreeRegressor()),
11    ("Random Forest", ensemble.RandomForestRegressor()),
12    ("AdaBoost", ensemble.AdaBoostRegressor()),
13    ("ExtraTree", ensemble.ExtraTreesRegressor()),
14    ("GradientBoosting", ensemble.GradientBoostingRegressor()),
15    ("K Neighbors", neighbors.KNeighborsRegressor()),
16 ]
17
18 model_names = []
19 accuracies = []
20
21 # Fitting models to Training Dataset and Scoring them on Test set
22 for dataset_name, dataset in [("Akosomba_Data", akosombo)]:
23     for model_name, model in models:
24         regressor_model = model
25         regressor_model.fit(X_train, y_train)
26
27         accuracy = regressor_model.score(X_test, y_test)
28         print(dataset_name, model_name, accuracy)
29
30         model_names.append(model_name)
31         accuracies.append(accuracy)

```

Akosomba_Data Linear Regression 0.9697494269312791
 Akosomba_Data Ridge 0.9697837468881153
 Akosomba_Data SVR 0.9696567718462318
 Akosomba_Data NuSVR 0.9761908456257131
 Akosomba_Data SVR 0.9762842898104876
 Akosomba_Data Decision Tree 0.9622653656060981
 Akosomba_Data Random Forest 0.9764322782401315
 Akosomba_Data AdaBoost 0.9569520639911224
 Akosomba_Data ExtraTree 0.9757105686833952
 Akosomba_Data GradientBoosting 0.9772234044519424
 Akosomba_Data K Neighbors 0.9742917147982414

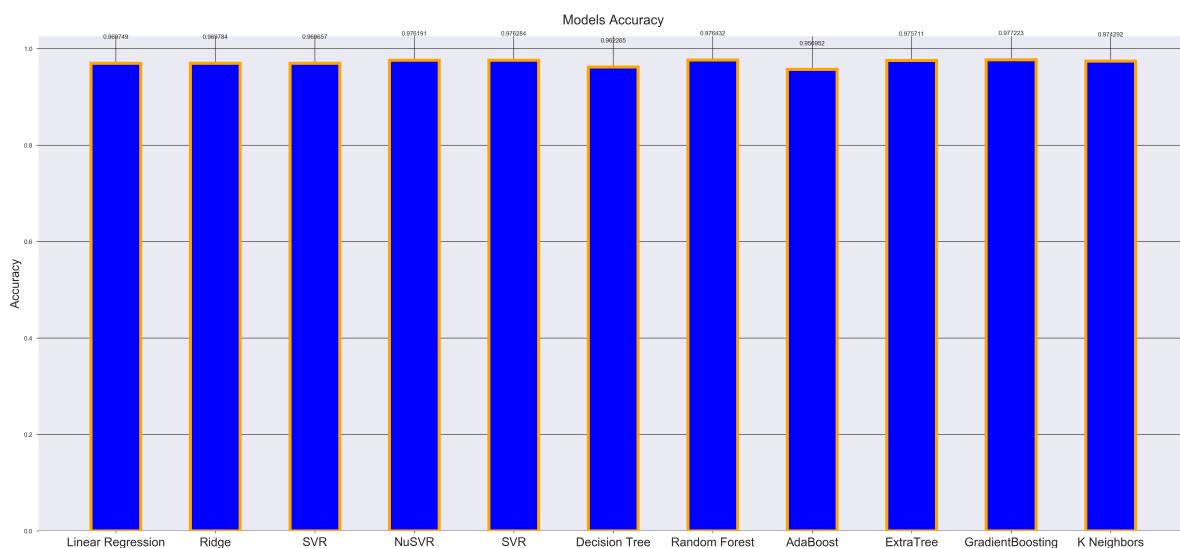
Visualizing Models' Accuracy with Bar Charts

In [7]:

```

1 # Size in inches (width, height) & resolution(DPI)
2 plt.figure(figsize=(34, 15), dpi=200)
3
4 x_loc = np.arange(len(models)) # the x locations for the groups
5 width = 0.5 # bar width
6
7 # plotting the graphs with bar chart
8 models_graph = plt.bar(
9     x_loc, accuracies, width, color="blue", edgecolor="orange", linewidth=5,
10 )
11
12 plt.title("Models Accuracy", fontsize=22, pad=20)
13 plt.xticks(x_loc, model_names, fontsize=20)
14 plt.ylabel("Accuracy", fontsize=20)
15 plt.grid(b=True, which="both", axis="both", color="black", linewidth=0.8)
16
17 # adding model accuracy on top of every bar
18 def addLabel(models):
19     for model in models:
20         height = model.get_height()
21         plt.text(
22             model.get_x() + model.get_width() / 2.0,
23             1.05 * height,
24             "%f" % height,
25             ha="center",
26             va="bottom",
27         )
28
29
30 addLabel(models_graph)
31
32 plt.savefig('Bar_Charts_of_Models_and_their_Accuracy.png', dpi=300, transparent=True)
33
34 plt.show()

```



Evaluating Models

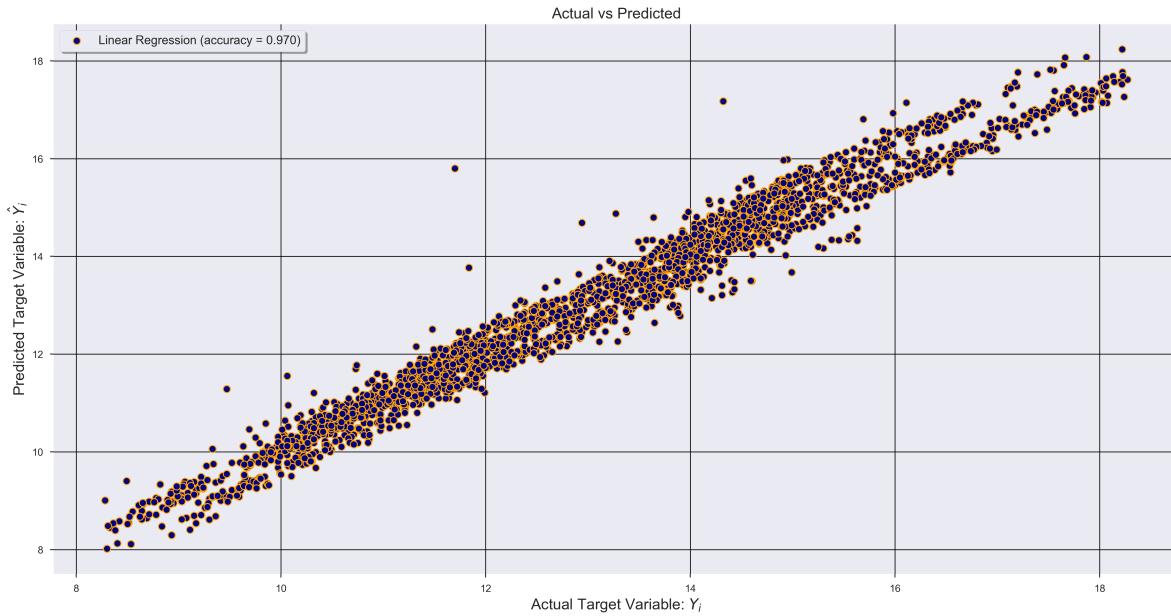
In [8]:

```

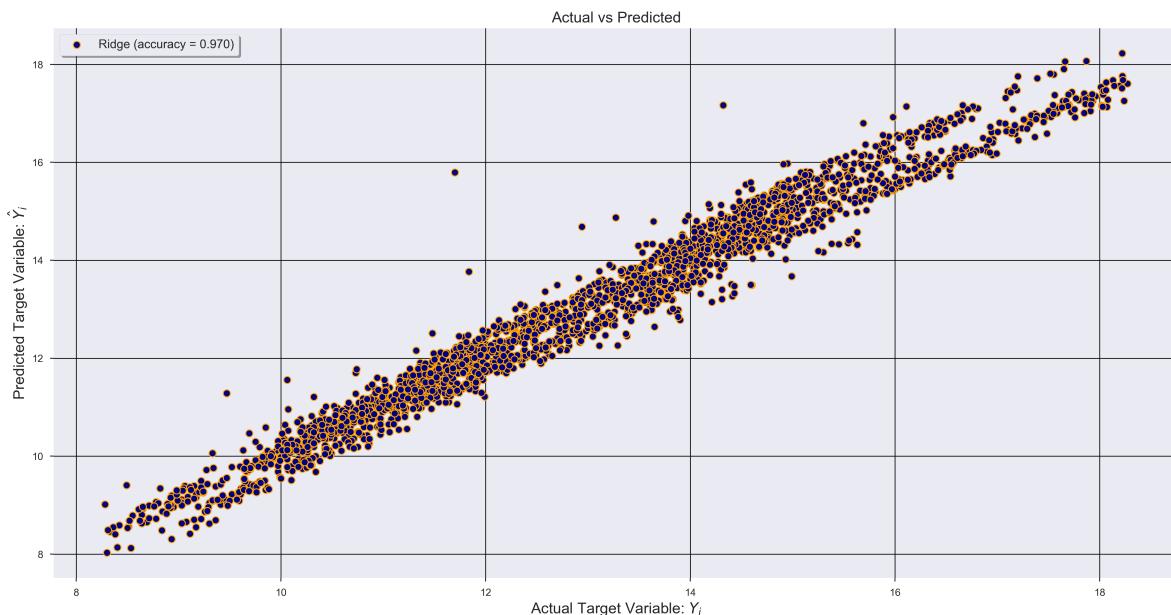
1 # Model Evaluation
2 for model_name, model in models:
3     helper.evaluate(X_test, y_test, model_name, model)

```

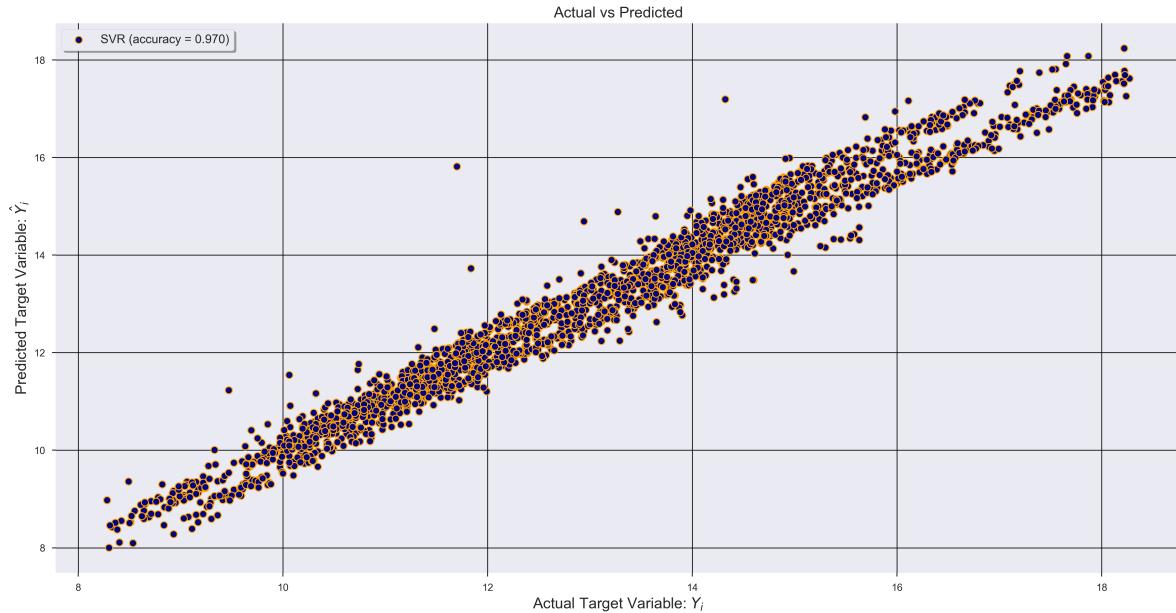
Linear Regression Mean Squared Error: 0.14625233556901102
 Linear Regression Root Mean Squared Error: 0.38242951712571954
 Linear Regression R2 Score: 0.969749426931279
 Linear Regression Explained Variance Score: 0.969772234576705
 Linear Regression Mean Absolute Error: 0.29263695465956613
 Linear Regression Meadian Abosolute Error: 0.2677789934798991
 Linear Regression Mean Squared Log Error: 0.0007141004739016008



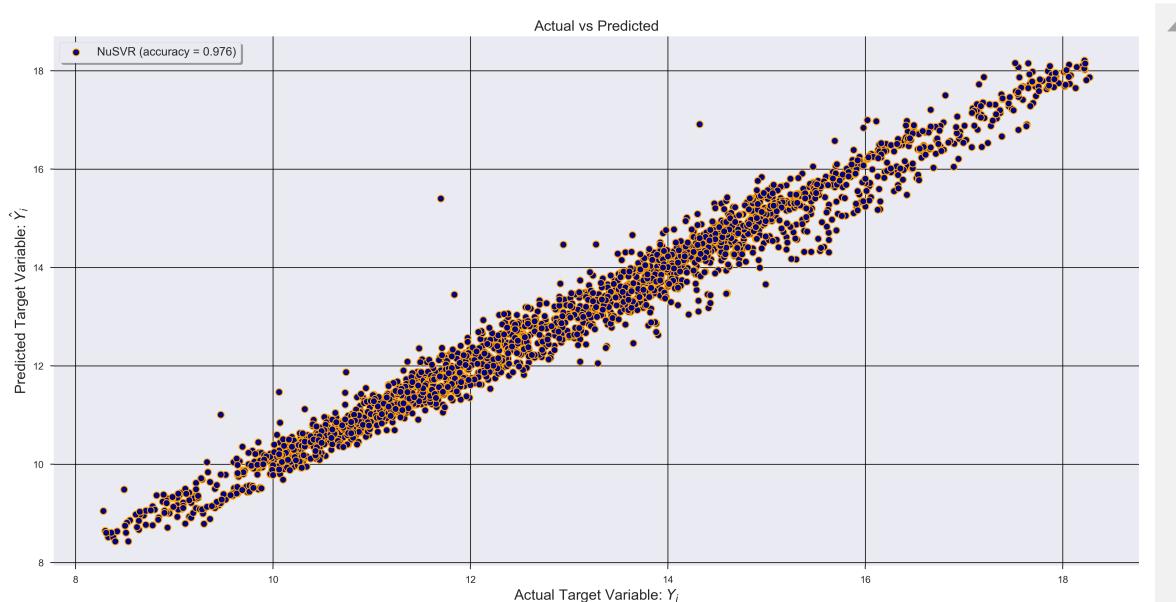
Ridge Mean Squared Error: 0.14608640899854475
 Ridge Root Mean Squared Error: 0.3822125181081132
 Ridge R2 Score: 0.9697837468881153
 Ridge Explained Variance Score: 0.9698066436041364
 Ridge Mean Absolute Error: 0.2926293335519423
 Ridge Meadian Abosolute Error: 0.26566972718727655
 Ridge Mean Squared Log Error: 0.0007123218460580986



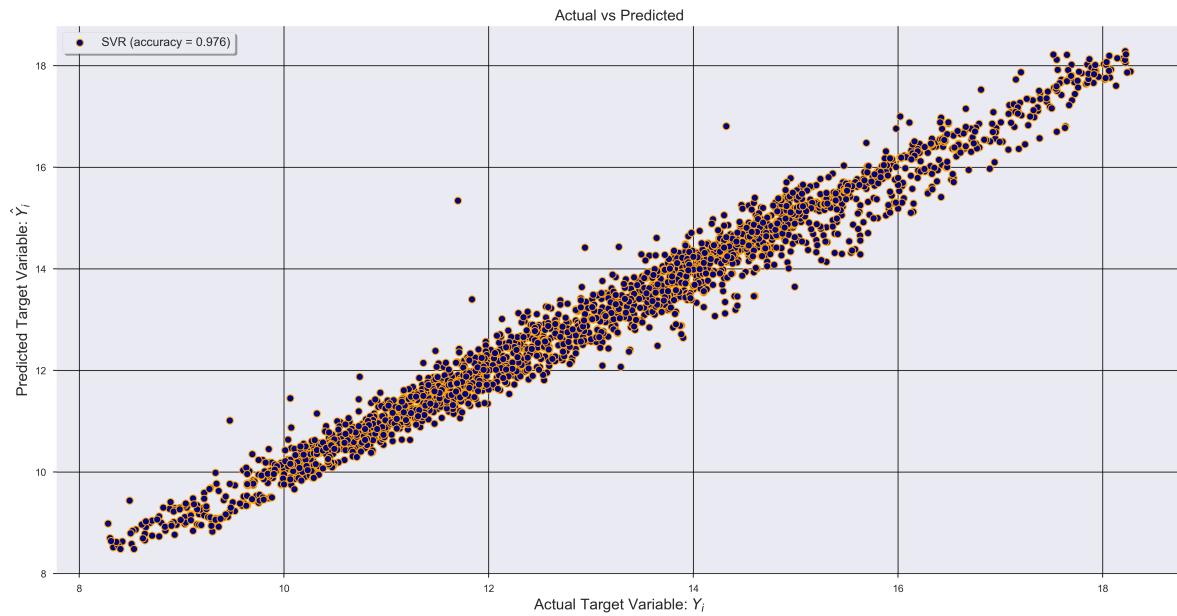
SVR Mean Squared Error: 0.14670029477162583
 SVR Root Mean Squared Error: 0.383014744848845
 SVR R2 Score: 0.9696567718462318
 SVR Explained Variance Score: 0.969658812471311
 SVR Mean Absolute Error: 0.2912638394608893
 SVR Median Abosolute Error: 0.2636070991213728
 SVR Mean Squared Log Error: 0.0007125399126440713



NuSVR Mean Squared Error: 0.1151100320397934
 NuSVR Root Mean Squared Error: 0.3392786937604444
 NuSVR R2 Score: 0.9761908456257131
 NuSVR Explained Variance Score: 0.9761920945378169
 NuSVR Mean Absolute Error: 0.24964070901547153
 NuSVR Median Abosolute Error: 0.1964009749299347
 NuSVR Mean Squared Log Error: 0.000575800094914912

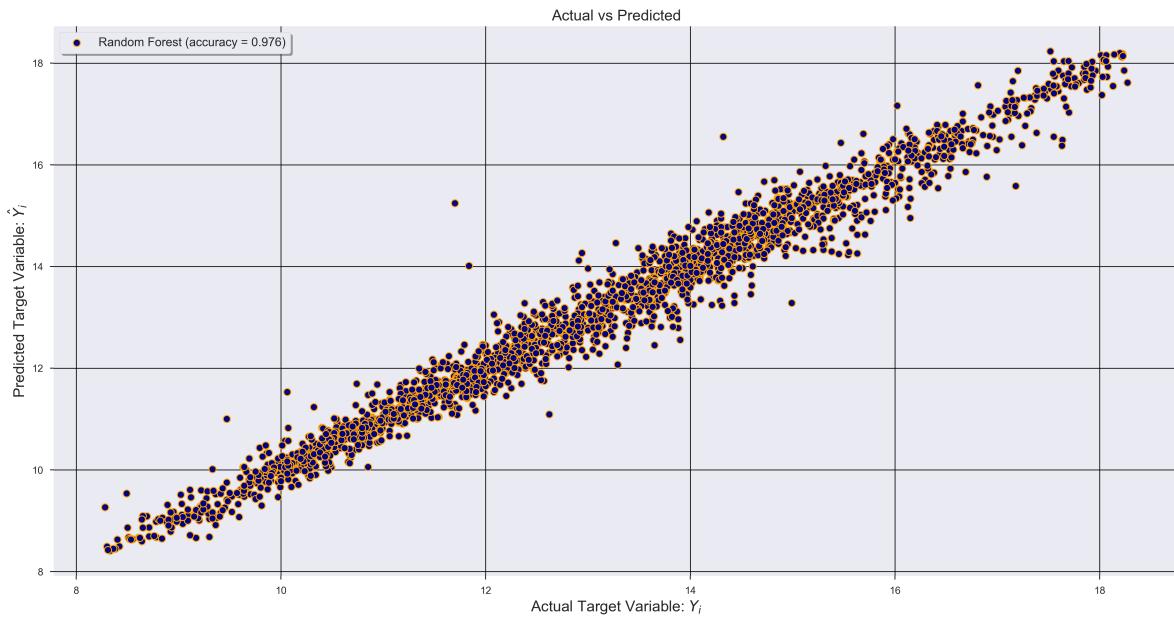


SVR Mean Squared Error: 0.1146582577795977
SVR Root Mean Squared Error: 0.338612252849181
SVR R2 Score: 0.9762842898104876
SVR Explained Variance Score: 0.9762930859445609
SVR Mean Absolute Error: 0.2442590031041782
SVR Meadian Abosolute Error: 0.18377625009494558
SVR Mean Squared Log Error: 0.000572986557719098

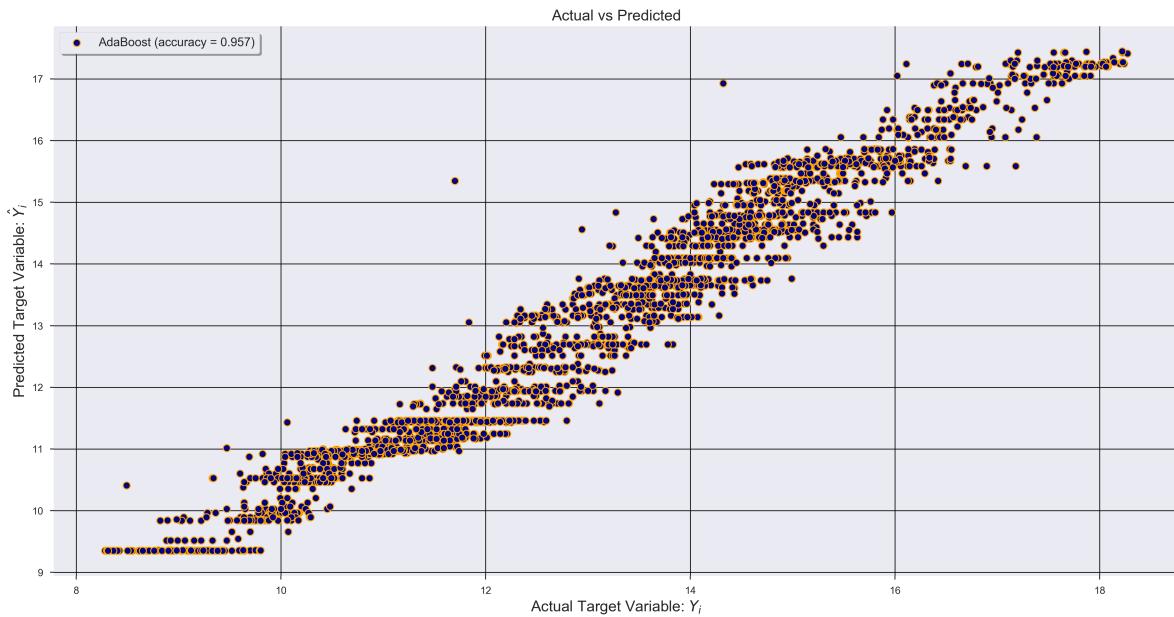


Decision Tree Mean Squared Error: 0.18243549963214703
Decision Tree Root Mean Squared Error: 0.4271246886239978
Decision Tree R2 Score: 0.9622653656060981
Decision Tree Explained Variance Score: 0.9623485228821076
Decision Tree Mean Absolute Error: 0.2804296681327466
Decision Tree Meadian Abosolute Error: 0.16000000000000014
Decision Tree Mean Squared Log Error: 0.0009007957414350518

Random Forest Mean Squared Error: 0.11394277865715507
 Random Forest Root Mean Squared Error: 0.3375541121911494
 Random Forest R2 Score: 0.9764322782401315
 Random Forest Explained Variance Score: 0.9764535491742466
 Random Forest Mean Absolute Error: 0.2311272736810036
 Random Forest Meadian Abosolute Error: 0.15338000000000207
 Random Forest Mean Squared Log Error: 0.0005630956765314753



AdaBoost Mean Squared Error: 0.20812369962120095
 AdaBoost Root Mean Squared Error: 0.4562057645637558
 AdaBoost R2 Score: 0.9569520639911224
 AdaBoost Explained Variance Score: 0.9569776365109394
 AdaBoost Mean Absolute Error: 0.3599828662011853
 AdaBoost Meadian Abosolute Error: 0.3025436447166907
 AdaBoost Mean Squared Log Error: 0.0011126190000955328



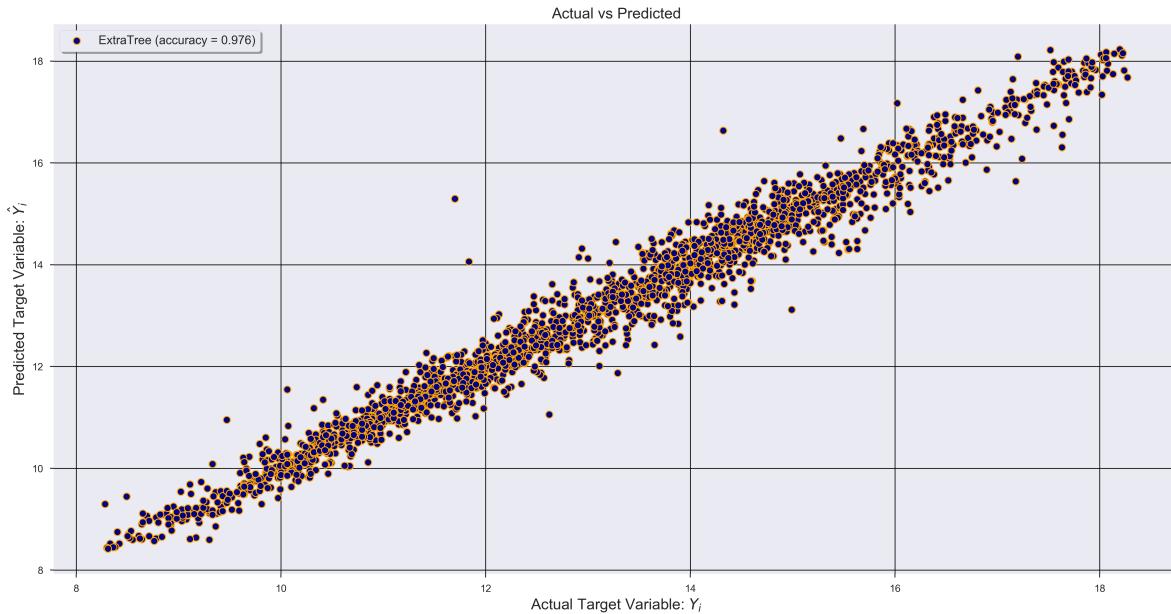
ExtraTree Mean Squared Error: 0.11743202522565402
 ExtraTree Root Mean Squared Error: 0.3426835642770952
 ExtraTree R2 Score: 0.9757105686833952

ExtraTree Explained Variance Score: 0.9757317960090854

ExtraTree Mean Absolute Error: 0.2303212874850074

ExtraTree Meadian Abosolute Error: 0.14989999999998638

ExtraTree Mean Squared Log Error: 0.0005826065648577845



GradientBoosting Mean Squared Error: 0.11011792363889294

GradientBoosting Root Mean Squared Error: 0.33184020799007

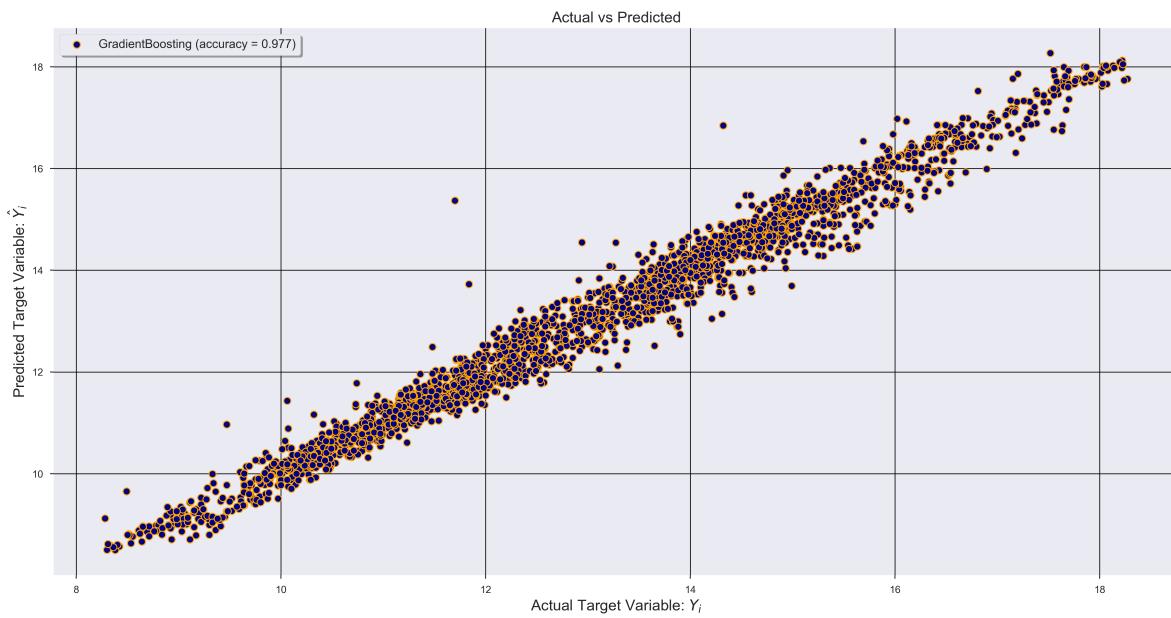
GradientBoosting R2 Score: 0.9772234044519423

GradientBoosting Explained Variance Score: 0.9772262384355532

GradientBoosting Mean Absolute Error: 0.24210040978092848

GradientBoosting Meadian Abosolute Error: 0.18397069087037288

GradientBoosting Mean Squared Log Error: 0.0005519263607088787



K Neighbors Mean Squared Error: 0.1242917529426631

K Neighbors Root Mean Squared Error: 0.3525503551872599

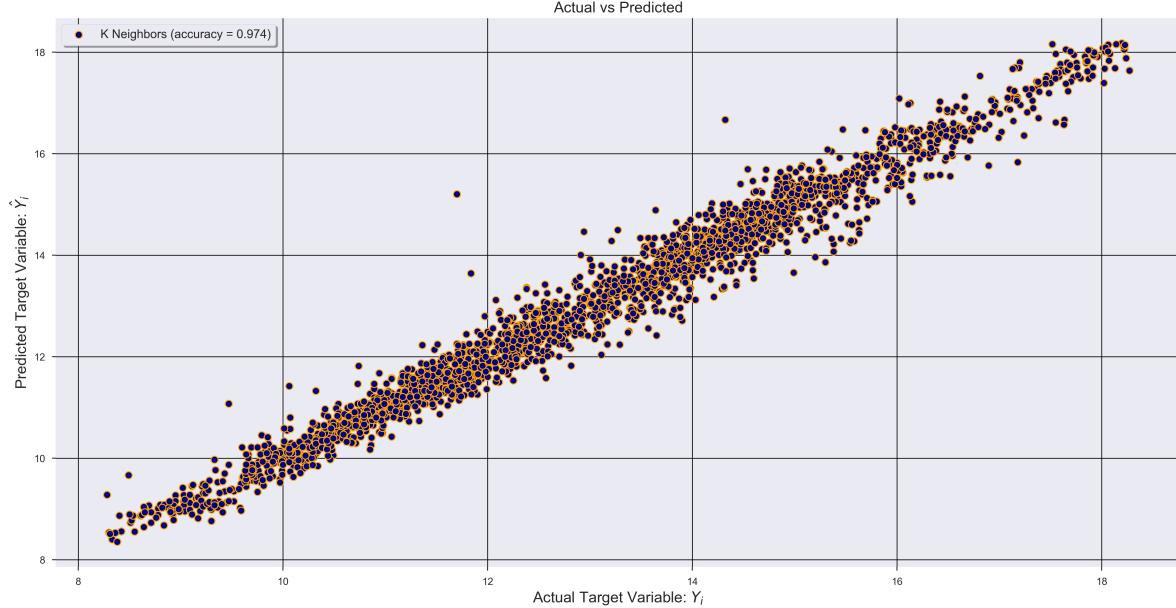
K Neighbors R2 Score: 0.9742917147982414

K Neighbors Explained Variance Score: 0.9743001723319055

K Neighbors Mean Absolute Error: 0.25122119152339034

K Neighbors Meadian Abosolute Error: 0.1819999999999986

K Neighbors Mean Squared Log Error: 0.0006201934569084686



Cross Validating Models

Cross Validating with a single metric

In [9]:

```

1 # Splitting data into 10 folds
2 cv_kfold = model_selection.KFold(n_splits=10, shuffle=True, random_state=23)
3 scorer = "r2"
4
5 model_names = []
6 cv_mean_scores = []
7 cv_std_scores = []
8
9 for model_name, model in models:
10     regressor_model = model
11     model_scores = model_selection.cross_val_score(
12         regressor_model, X, y, cv=cv_kfold, scoring=scorer, n_jobs=-1, verbose=1,
13     )
14
15     print(
16         f"{model_name} Accuracy: {0.2f} (+/- {0.2f})"
17         % (model_scores.mean(), model_scores.std() * 2)
18     )
19
20     model_names.append(model_name)
21     cv_mean_scores.append(model_scores.mean())
22     cv_std_scores.append(model_scores.std())

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 6 out of 10 | elapsed: 1.1s remaining:
0.7s

[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed: 1.1s finished

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 6 out of 10 | elapsed: 0.0s remaining:
0.0s

[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed: 0.0s finished

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

Linear Regression Accuracy: 0.97 (+/- 0.00)

Ridge Accuracy: 0.97 (+/- 0.00)

[Parallel(n_jobs=-1)]: Done 6 out of 10 | elapsed: 1.9s remaining:
1.3s

[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed: 2.5s finished

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

SVR Accuracy: 0.86 (+/- 0.21)

[Parallel(n_jobs=-1)]: Done 6 out of 10 | elapsed: 9.0s remaining:
6.0s

[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed: 14.4s finished

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

NuSVR Accuracy: 0.93 (+/- 0.00)

[Parallel(n_jobs=-1)]: Done 6 out of 10 | elapsed: 7.9s remaining:
5.2s

[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed: 12.7s finished

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 6 out of 10 | elapsed: 0.0s remaining:
0.0s

[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed: 0.0s finished

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

SVR Accuracy: 0.94 (+/- 0.00)

```
Decision Tree Accuracy: 0.95 (+/- 0.00)
[Parallel(n_jobs=-1)]: Done  6 out of 10 | elapsed:  3.5s remaining:
2.3s
[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed:  6.0s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

Random Forest Accuracy: 0.98 (+/- 0.00)

```
[Parallel(n_jobs=-1)]: Done  6 out of 10 | elapsed:  0.6s remaining:
0.3s
[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed:  1.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

AdaBoost Accuracy: 0.95 (+/- 0.00)

```
[Parallel(n_jobs=-1)]: Done  6 out of 10 | elapsed:  2.5s remaining:
1.7s
[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed:  4.1s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

ExtraTree Accuracy: 0.98 (+/- 0.00)

```
[Parallel(n_jobs=-1)]: Done  6 out of 10 | elapsed:  1.1s remaining:
0.7s
```

GradientBoosting Accuracy: 0.98 (+/- 0.00)

K Neighbors Accuracy: 0.98 (+/- 0.00)

```
[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed:  1.7s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done  6 out of 10 | elapsed:  0.0s remaining:
0.0s
[Parallel(n_jobs=-1)]: Done 10 out of 10 | elapsed:  0.0s finished
```

In [10]:

```
1 cv_results = pd.DataFrame({"model_name": model_names, "mean_score": cv_mean_scores, "std_score": cv_std_scores})
2 cv_results.sort_values("mean_score", ascending=False, inplace=True)
3 cv_results.to_csv("cross_validation_results.csv", index=True)
4 cv_results
```

Out[10]:

	model_name	mean_score	std_score
6	Random Forest	0.978089	0.001374
9	GradientBoosting	0.978030	0.001234
8	ExtraTree	0.977338	0.001383
10	K Neighbors	0.975400	0.001267
1	Ridge	0.971263	0.001156
0	Linear Regression	0.971263	0.001156
5	Decision Tree	0.964602	0.002421
7	AdaBoost	0.952748	0.002287
4	SVR	0.941600	0.001811
3	NuSVR	0.934291	0.002104
2	SVR	0.857605	0.103771

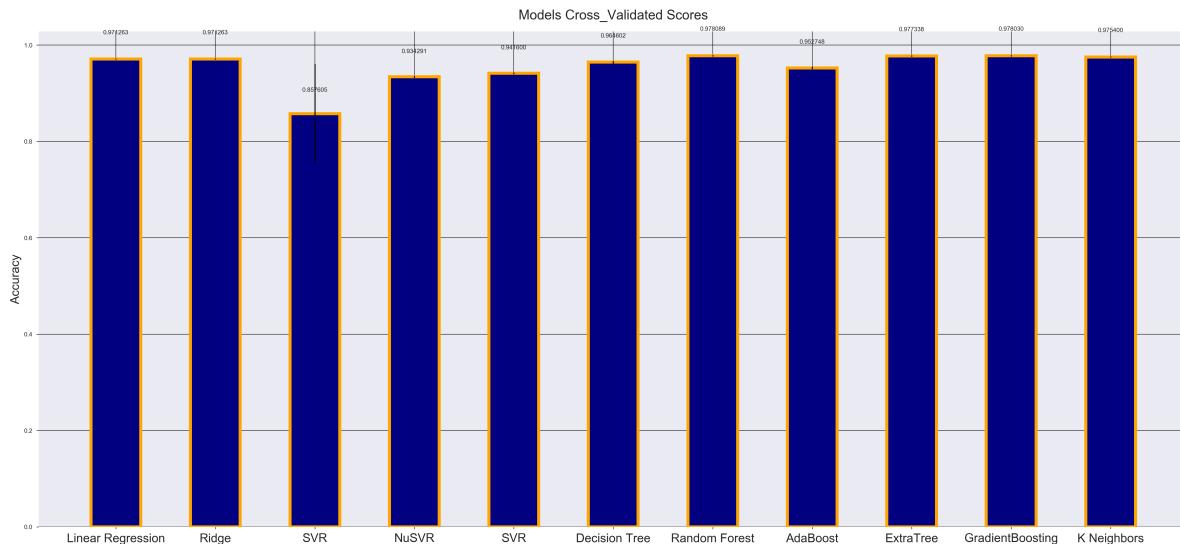
Visualizing Cross Validated Models with Bar Charts

In [11]:

```

1 plt.figure(figsize=(34, 15), dpi=200)
2
3 x_loc = np.arange(len(models))
4 width = 0.5
5
6 models_graph = plt.bar(
7     x_loc, cv_mean_scores, width, yerr=cv_std_scores, color="navy", edgecolor="orange",
8 )
9 plt.title("Models Cross_Validated Scores", fontsize=22, pad=20)
10 plt.xticks(x_loc, model_names, fontsize=20)
11 plt.ylabel("Accuracy", fontsize=20)
12 plt.grid(b=True, which="both", axis="both", color="black", linewidth=0.8)
13
14 addLabel(models_graph)
15
16 plt.savefig('Bar_Charts_of_Cross_Validated_Models_and_their_Accuracy.png', dpi=300, tra...
17
18 plt.show()

```



Training the Model with the Highest Score with Default Hyperparameters

In [12]:

```

1 # Instantiating model object
2 high_score_model = ensemble.GradientBoostingRegressor()
3
4 # Fitting the model on Train set
5 high_score_model.fit(X_train, y_train)
6
7 # Scoring the model on Test set
8 high_score_model_accuracy = high_score_model.score(X_test, y_test)
9
10 print(
11     f"Model without tuned hyperparameters has an accuracy of {high_score_model_accuracy}
12 )

```

Model without tuned hyperparameters has an accuracy of 0.9772167213704891

Predicting with the Trained Model and Saving Predicted Results as csv

In [13]:

```
1 y_pred = high_score_model.predict(X_test)
2
3 data = pd.DataFrame({"generation": list(y_test), "predicted_generation": list(y_pred),
4 data.to_csv("model_predicted_values.csv", index=True)}
```

In [14]:

```
1 data.head(10)
```

Out[14]:

	generation	predicted_generation
0	17.596	17.663987
1	15.630	15.226165
2	10.850	11.354292
3	14.520	14.603583
4	8.420	8.569025
5	13.640	14.509512
6	14.010	13.864608
7	9.699	9.489637
8	13.900	12.741841
9	11.150	11.255741

Optimizing the Hyperparameter of the Best Model with GridSearchCV

In [15]:

```

1 # Kfold with n_splits = 5 to split the Dataset into 5-folds
2 kfold = model_selection.KFold(n_splits=5, shuffle=True, random_state=23)
3
4 # Dictionary of parameters to tune
5 parameters = {
6     "loss" : ['ls', 'lad', 'huber', 'quantile'],
7     "n_estimators" : [120, 500, 800, 1200],
8     "max_depth" : [15, 25, 30, None],
9     "min_samples_split" : [5, 10, 15, 100],
10    "min_samples_leaf" : [1, 2, 5, 10],
11    "max_features" : ["log2", "sqrt", None],
12 }
13
14 scorer = "r2"
15
16 # Instantiating Search object
17 grid = model_selection.RandomizedSearchCV(
18     estimator=high_score_model,
19     param_distributions=parameters,
20     scoring=scorer,
21     cv=kfold,
22     n_jobs=-1,
23     verbose=1,
24 )
25
26 # Fit the grid object on Training Dataset
27 grid.fit(X_train, y_train)

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 10.0min
[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 11.2min finished

Out[15]:

```

RandomizedSearchCV(cv=KFold(n_splits=5, random_state=23, shuffle=True),
                    error_score=nan,
                    estimator=GradientBoostingRegressor(alpha=0.9, ccp_alpha=
a=0.0,
                                                criterion='friedman
_mse',
                                                init=None,
                                                learning_rate=0.1,
                                                loss='ls', max_dept
h=3,
                                                max_features=None,
                                                max_leaf_nodes=Non
e,
                                                min_impurity_decrea
se=0.0,
                                                min_impurity_split=
None,
                                                min_samples_leaf=1,
                                                min_samples_split=
2,
                                                min...
                                                iid='deprecated', n_iter=10, n_jobs=-1,
param_distributions={'loss': ['ls', 'lad', 'huber',

```

```

        'quantile'],
        'max_depth': [15, 25, 30, None],
        'max_features': ['log2', 'sqrt', N
one],
        'min_samples_leaf': [1, 2, 5, 10],
        'min_samples_split': [5, 10, 15, 1
00],
        'n_estimators': [120, 500, 800, 12
00]},
    pre_dispatch='2*n_jobs', random_state=None, refit=True,
    return_train_score=False, scoring='r2', verbose=1)

```

In [16]:

```

1 results = pd.DataFrame(grid.cv_results_)[["params", "mean_test_score", "std_test_score"]
2 results.sort_values("rank_test_score", inplace=True)
3 results.to_csv("hyperparameter_tuning_results.csv", index=True)
4 results

```

Out[16]:

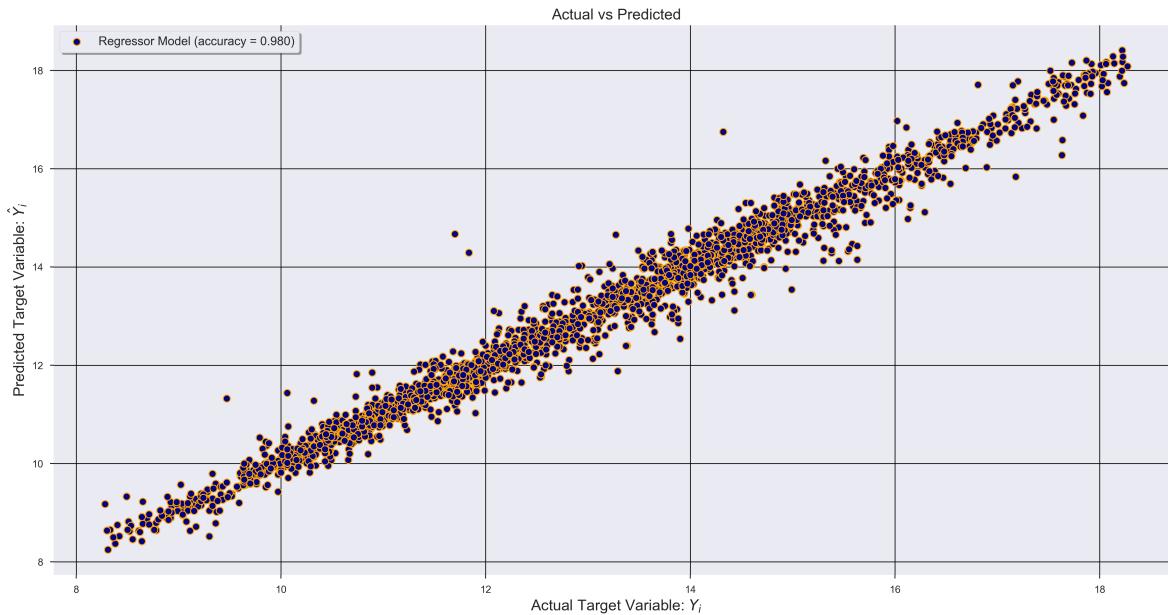
	params	mean_test_score	std_test_score	rank_test_score
9	{'n_estimators': 500, 'min_samples_split': 100...}	0.980954	0.001212	1
6	{'n_estimators': 800, 'min_samples_split': 100...}	0.980632	0.000622	2
7	{'n_estimators': 500, 'min_samples_split': 100...}	0.979432	0.000919	3
4	{'n_estimators': 120, 'min_samples_split': 10,...}	0.978602	0.000493	4
2	{'n_estimators': 120, 'min_samples_split': 5, ...}	0.978561	0.000687	5
0	{'n_estimators': 1200, 'min_samples_split': 10...}	0.977948	0.000639	6
3	{'n_estimators': 1200, 'min_samples_split': 10...}	0.977863	0.000600	7
5	{'n_estimators': 800, 'min_samples_split': 10,...}	0.976270	0.000749	8
8	{'n_estimators': 800, 'min_samples_split': 5, ...}	0.974772	0.000898	9
1	{'n_estimators': 1200, 'min_samples_split': 10...}	0.968720	0.001231	10

Evaluating the Best Estimator from the GridSearch

In [17]:

```
1 best_estimator = grid.best_estimator_
2 helper.evaluate(X_test, y_test, "Regressor Model", best_estimator)
```

Regressor Model Mean Squared Error: 0.09690800378247133
Regressor Model Root Mean Squared Error: 0.31130050398685727
Regressor Model R2 Score: 0.9799557207892775
Regressor Model Explained Variance Score: 0.9799816608782863
Regressor Model Mean Absolute Error: 0.20714495307769037
Regressor Model Median Absolute Error: 0.13735171170935345
Regressor Model Mean Squared Log Error: 0.0004856029538959505



Predicting with the Best Estimator and Saving Predicted Results as csv

In [18]:

```

1 tune_y_pred = best_estimator.predict(X_test)
2
3 hyp_tune_data = pd.DataFrame(
4     {"generation": list(y_test), "predicted_generation": list(tune_y_pred),}
5 )
6
7 hyp_tune_data.to_csv("tune_model_predicted_values.csv", index=True)
8
9 hyp_tune_data.head(10)

```

Out[18]:

	generation	predicted_generation
0	17.596	17.759893
1	15.630	15.338581
2	10.850	11.194873
3	14.520	14.460236
4	8.420	8.525759
5	13.640	14.225518
6	14.010	13.788593
7	9.699	9.588924
8	13.900	12.536513
9	11.150	10.931286

Saving the Best Estimator with joblib

In [19]:

```

1 import joblib
2
3 joblib.dump(best_estimator, "GradientBoosting.joblib")

```

Out[19]:

```
['GradientBoosting.joblib']
```

Feature Importance

In [20]:

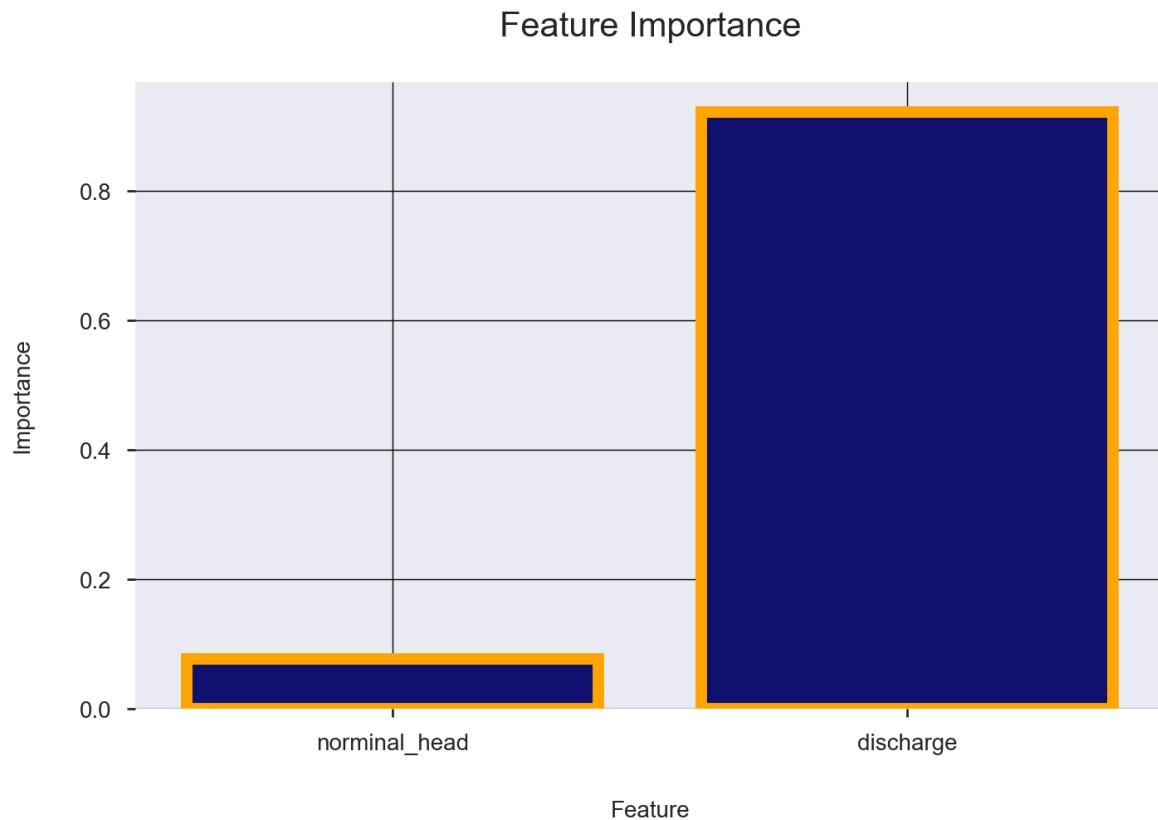
```
1 features = ['norminal_head', 'discharge']
2 importances = best_estimator.feature_importances_
3
4 feature_importance = pd.DataFrame({
5     'feature': features,
6     'importance': importances,
7 })
8
9 feature_importance.sort_values("importance", inplace=True, ascending=False)
10 feature_importance.to_csv("feature_importance.csv", index=True)
11 feature_importance
```

Out[20]:

	feature	importance
1	discharge	0.92209
0	norminal_head	0.07791

In [21]:

```
1 import seaborn as sns
2
3 plt.figure(figsize=(8, 5), dpi=200)
4
5 sns.barplot(x=features, y=importances, color="navy", edgecolor="orange", linewidth=5)
6 plt.title("Feature Importance", size=15, pad=20)
7 plt.xlabel("Feature", fontsize=10, labelpad=20)
8 plt.ylabel("Importance", fontsize=10, labelpad=20)
9
10 plt.grid(b=True, which="both", axis="both", color="black", linewidth=0.5)
11
12 plt.savefig('feature_importance.png', dpi=300, transparent=True)
13
14 plt.show()
```



In []:

1

In []:

1